

Neural Network Language Models

Natural Language Processing

Daniel Ortiz Martínez

Table of Contents

1. Motivation and Foundations
2. Recurrent Models
3. Convolutional Models
4. Sequence to Sequence Models
5. Transformer-based Models
6. The EuTrans Dataset
7. The Era of Foundation Models

Motivation and Foundations

- Language models (LMs) calculate the probability of the next word in a sequence given the preceding ones
- One popular LM implementation are the so-called N-gram LMs
- N-gram LMs very successful in many NLP tasks
- N-gram LMs define a probability distribution for each word given the preceding ones:


$$p(x|h) = \frac{C(h, x)}{C(h)}$$

Limitations of N-gram LMs

- Many histories are similar but exact match of h is assumed
- Arbitrary representation of words
- Poor at handling uncommon or unseen events
- Poor at capturing long term relationships between words
- Curse of dimensionality

Neural Networks

- Neural networks (NNs) are statistical learning algorithms that estimate functions from a set of inputs
- NNs composed of connected neurons usually grouped in layers
- The canonical example is the feedforward NN:



- Neurons are processing units computing weighted sums of their inputs
 - Each connection has an associated weight
 - Output is given by an activation function over the weighted sum
- Three parameters define a specific NN:
 - Interconnection pattern between the neurons
 - Learning process used to update or train the weights
 - Activation function converting the neuron input into its output

- NNs can be used to build LMs, removing or alleviating the limitations of N-gram LMs
- The history h is projected into some continuous low-dimensional space, where similar histories get clustered
- Thanks to parameter sharing among similar histories, the model is more robust: less parameters have to be estimated


- Early NN-based algorithms such as the multilayer perceptron were typically composed of at most 3 layers
- Each layer can be seen as a mathematical representation of the input
- Multiple layers allow to capture the different factors or features that are relevant for a particular learning task
 - Opposed to hand-made features

- Deep learning models are NNs composed of many layers
- Deep learning has been made possible due to some advances:
 - Availability of large amounts of training data
 - More powerful computer infrastructures
 - Better optimization algorithms
- State-of-the-art LMs follow a deep architecture: transformer models

Recurrent Models

Recurrent Neural Networks


- RNNs are just networks with loops
- Loops allow the algorithm to capture long-range dependencies



$$h_t = \tanh(W_h h_{t-1} + W_x x_t + b)$$

Recurrent Neural Networks

- It's typical to represent RNNs unrolled in time:



- In simple RNNs input is processed in one direction
- Bidirectional RNNs concatenates two simple RNNs in both directions


Specializations of the basic RNN

- Initial implementations of the recurrent unit (A) were very simple, summing 2 matrix multiplications and applying the activation function
- This configuration was not able to preserve information over many timesteps due to the vanishing gradients problem
- To avoid this, alternative recurrent units have been defined:
 - **GRU** (Gated Recurrent Unit): use reset and update gates to determine the information to be retained for future predictions (Cho et al. 2014)
 - **LSTM** (Long Short Term Memory): similar to GRU units but using three gates instead of two (Hochreiter and Schmidhuber 1997)

The Vanishing Gradient Problem

How does an RNN learn?

- After a prediction error, the network sends a **correction signal** backwards through time to adjust its weights. This is called *back-propagation through time* (BPTT)
- Each step backwards, the signal is multiplied by a number smaller than 1
- Over many steps, the signal shrinks exponentially



The Vanishing Gradient Problem: Consequences

The RNN have difficulties to learn

- Long-range agreement:
*"The keys on the table **are** missing"* (subject and verb are far apart)
- Coreference across sentences:
*"Mary told Ann that **she** had won"* (who is *she*?)
- In general: any dependency spanning more than a few tokens is unreliable


The Vanishing Gradient Problem: LSTM & GRU

Both architectures add gating mechanisms, learnable switches that decide:


Gate	Role
Forget	What to erase from memory
Input	What new information should be stored
Output	What to pass to the next step

Crucially, information can flow through a dedicated cell state (LSTM) or update gate (GRU) without being repeatedly multiplied down, giving the network a stable path to carry context over long distances

- RNNs can be used to implement LMs ([Mikolov et al. 2010](#))



RNN Language Models: Architecture



RNN Language Models: Architecture

- **Embedding layer:** maps categorical variable to a continuous one
- **Linear layer:** maps RNN output to an alphabet size vector
- **Softmax layer:** converts output to probability distribution


$$x_t = E(w_t)$$

$$h_t = \tanh(W_h h_{t-1} + W_x x_t + b)$$

$$z_t = W_y h_t + b_y$$

$$P(w_{t+1} \mid w_1, \dots, w_t) = \text{softmax}(z_t)$$


RNN Language Models: Architecture



Bidirectional RNN Language Models

- RNN LMs only take into account left context when predicting next symbol
- Bidirectional RNN LMs enables use of both right and left contexts
- Used for tasks where the full sequence is available at inference time (text classification, sequence labeling, etc.)


Bidirectional RNN Language Models



Convolutional Models

Convolutional Neural Networks for Image Analysis

- Convolutional neural networks (CNNs) are a class of NNs originally used for image analysis
- Biologically-inspired connection pattern (visual cortex)




Convolutional Neural Networks for Image Analysis

- Convolutions are learned filters that identify features in the input
- A limitation of this approach is that it records the precise position of the features in the input
- One way to alleviate this problem is subsampling, where a lower resolution version of the input is created
- Subsampling is typically implemented by means of a pooling layer

Convolutional Neural Networks for Text Classification

- CNNs can also be applied to NLP, although initially they were not used for language modelling
- One important application of CNNs in NLP would be text classification
- Seminal work applied convolution and pooling operations (Kim 2014)


Convolutional Neural Networks for Text Classification




Convolutional Neural Networks for Language Modeling

- CNNs for language modeling require some adaptations:
 - Pooling is not executed so as to keep positional information
 - To predict a new word it's necessary to ensure that only the previous ones are considered (padding symbols are used when necessary)
 - Context size is fixed, but it can be very long when stacking many layers
 - Stacking many layers make training difficult, requiring the use of specific techniques such as residual connections
- See for example ([Dauphin et al. 2016](#)) for more information

Convolutional Neural Networks for Language Modeling



Convolutional Neural Networks for Language Modeling



Sequence to Sequence Models

Conditional Language Models

- Conventional language models estimate the probability $p(y)$ of a sequence of tokens $(y_1, y_2, \dots, y_n)$
- In conditional language models, the unconditional probability $p(y)$ is replaced by a conditional one:

$$p(y|x) = p(y_1, y_2, \dots, y_n|x)$$

- x can be defined in multiple ways (text, image, acoustic signal)
- Sequence-to-sequence (Seq-to-Seq) models estimate $p(y|x)$


- Seq-to-Seq models (Sutskever et al. 2014) were initially defined to solve machine translation (MT) tasks
- For a given input sentence x , the statistical approach to MT finds the translation of highest probability in the output language, y

$$\hat{y} = \arg \max_y \{P(y|x)\} = \arg \max_y \{P(y) \cdot P(x|y)\}$$

- Original statistical MT systems modeled the distributions after applying the Bayes rule:
 - $P(y)$ modeled with an N-gram language model
 - $P(x|y)$ modeled with an HMM-based alignment model


Machine Translation: Alignments

- Input to output alignment is a key idea in MT
- It can be graphically represented as a matrix
- HMM-based MT technology generates alignments as a by-product




The Encoder-Decoder Framework

- The encoder-decoder framework is the standard modeling paradigm for Seq-to-Seq tasks



Seq-to-Seq Models

- Original Seq-to-Seq model concatenates two RNN LMs ([Sutskever et al. 2014](#))
- First RNN LM encodes input sequence into a vector
- Second RNN LM decodes this vector into the output sequence




Seq-to-Seq Plus Attention Models


- Original Seq-to-Seq models had problems to capture long range dependencies between symbols
 - the signal should propagate through many time steps in the encoder
 - the decoder should work with just one representation of the input
 - difficult to reflect that each output token can be related to different parts of the input
- The solution is the attention mechanism ([Bahdanau et al. 2015](#))
 - when generating a new output token, the decoder is granted access to all of the encoder states
 - this capability is used to decide which input parts are more relevant

Seq-to-Seq Attention Mechanism

- The Seq-to-Seq architecture is modified to include an attention layer





Seq-to-Seq Attention Mechanism (Detailed)



Seq-to-Seq Attention Mechanism

- The attention mechanism computes the relevance of the encoder states for each output token as a distribution
- If this information is represented graphically, we get something similar to the MT alignments we've seen before




Transformer-based Models

- RNN training is difficult to parallelize
- CNNs fast to train but not naturally designed to handle symbol dependencies
- How could we define a model that combines the best of both worlds?


Transformer Models

- Transformer models ([Vaswani et al. 2017](#)) are based solely on attention mechanisms (no recurrence or convolutions)
- Its definition enables huge training speed gains with respect to RNNs
- Transformer models constitute the state-of-the-art in language and Seq-to-Seq modelling
- Transformer models also adopt the encoder-decoder framework


Transformer Models



Transformer Models



Transformer Models



- **Encoder Self Attention:** all encoder states are received at once and this module establishes relationships between them
- **Decoder Self Attention:** basically the same as encoder attention but the output tokens are generated one at a time. During training masking is used to forbid the decoder to look ahead
- **Encoder-Decoder Attention:** this module establishes relationships between encoder and decoder states

Transformer: Scaled Dot-Product Attention

The scaled dot-product attention generates a contextual representation for each token in a sequence:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}}\right) \mathbf{V}$$

The three inputs are learned linear projections of the token representations (static embeddings), that are used to compare tokens one-to-one:

- $\mathbf{Q}$ (Queries) The embeddings of the tokens to be evaluated
- $\mathbf{K}$ (Keys) Again the token embeddings as keys of a key-value dictionary
- $\mathbf{V}$ (Values) The output embeddings of the tokens

$\sqrt{d_k}$ prevents the dot products from growing too large (and softmax from becoming too peaked) as d_k increases^a

^a d_k is the dimension of $\mathbf{Q}$ and $\mathbf{K}$ (a model hyperparameter).

Transformer: Scaled Dot-Product Attention

Calculation

1. **Score** every token against every other: QK^T gives a matrix of raw similarity scores
2. **Scale and normalize** with softmax to obtain attention weights $\alpha_{ij} \in [0, 1]$ summing to 1 per row
3. **Aggregate** the value vectors V weighted by α_{ij} : each output is a context-aware blend of all positions

Encoder: Q , K , V from the same sequence (self-attention)

Decoder: masked self-attention + cross-attention (K , V from encoder)


- Explainability of NN model's output is an open research problem
- The attention mechanism can be useful for this purpose
- RNN Seq-to-Seq attention mechanism produces one attention matrix
- Transformer models produce multiple attention matrices
- There are works that try to interpret this information (Voita et al. 2019)

- The data required to train a model for a particular task is often scarce
- However, sometimes there is plenty of data for a related task
- Transfer learning techniques try to transfer knowledge from one task to another
- Most relevant technique for this presentation is *model pre-training*
 - a model trained for a general task is *fine-tuned* for a more specific or *downstream task*
 - used in two popular models: GPT and BERT

Generative Pre-Training for Language Understanding

- Generative Pre-Training (GPT) ([Radford et al. 2018](#)) is a left-to-right language model
- GPT uses a 12-layer transformer decoder with no decoder-encoder attention
- At the pre-training stage, the model is trained to predict the next token of a sequence
- At the fine-tuning stage, the model allows to perform more specific tasks (e.g. sentence classification)

Generative Pre-Training for Language Understanding




NOTE: Figure taken from GPT's original paper

Bidirectional Encoder Representations from Transformers

- Bidirectional Encoder Representations from Transformers (BERT) ([Devlin et al. 2019](#)) is a bidirectional language model
- BERT is a transformer's encoder
- At the pre-training stage, BERT learns two tasks:
 - predict masked tokens in a sequence
 - determine whether a sequence of tokens goes after another one: next sentence prediction (NSP)
- At the fine-tuning stage, the model can be specialized to carry out tasks similar to those solved by GPT

Bidirectional Encoder Representations from Transformers



NOTE: Figure taken from BERT's original paper

The EuTrans Dataset

The EuTrans / Traveler Task Dataset

The EuTrans Dataset ([Amengual et al. 2001](#)) is a **parallel corpus** in the travel domain, developed within the European project *EuTrans* (Example-based and Statistical Machine Translation, 1996–1999)

- One of the first standard benchmarks for statistical MT
- Covers everyday tourist situations: hotels, restaurants, sightseeing
- Originally released as a Spanish–English parallel corpus
- Deliberately small and constrained to facilitate controlled experimentation

Corpus Statistics

	Spanish	English
Vocabulary size	~686 words	~513 words
Training sentences	10,000	10,000
Test sentences	2,996	2,996
Avg. sentence length	~7 tokens	~7 tokens

- Very limited vocabulary and short sentences
- High regularity and redundancy $\Rightarrow$ ideal for early MT models

Restricted domain: tourist conversations in three main scenarios

Hotel

Room reservations,
check-in/out, complaints,
billing

Restaurant

Ordering food and drinks,
asking for the menu, paying
the bill

Tourism

Asking for directions,
visiting monuments,
transport inquiries

Sentence Pair Example

ES: “¿Puede darme una habitación individual para esta noche?”

EN: “Can you give me a single room for tonight?”

EuTrans was used to evaluate a wide range of MT approaches:

- **IBM Models (1–5)**: word-level alignment models ([P. F. Brown et al. 1993](#))
- **Hidden Markov Models (HMM)**: alignment based on first-order Markov dependencies
- **Finite-State Transducers (FST)**: stochastic transducers exploiting the regularity of the domain
- **Example-Based MT (EBMT)**: translation by analogy with stored sentence pairs
- **Syntax-based models**: early hierarchical approaches

Legacy and Significance

- Established the practice of standardized corpus-based evaluation in MT research
- Served as a proving ground for probabilistic alignment models that underpin modern NMT
- Motivated the creation of larger, multi-domain corpora such as Europarl ([Koehn 2005](#)) and later WMT datasets
- Still cited in historical reviews and tutorial presentations on the evolution of MT

The Era of Foundation Models

Foundation Models

- GPT and BERT are examples of an emerging paradigm in artificial intelligence (AI): the so-called *foundation models* (Bommasani et al. 2021)
- A foundation model is a model trained on broad data that can be adapted to a wide range of downstream tasks
- Foundation models are based on deep learning and transfer learning
- Foundation models are closely linked to the concepts of *homogeneization* and *emergence*

Foundation Models: Homogeneization

- The introduction of BERT marks the beginning of the era of foundation models, characterized by an unprecedented level of homogeneization
- Homogeneization in AI refers to the capability to build machine learning systems that can work in a wide range of applications
- In NLP, almost all of the state-of-the-art models are adapted from one of a few foundation models, such as BERT
- This phenomenon is also being observed across research communities, where transformer-based sequence modeling is being applied to images, tabular data, protein sequences, etc.

- Another distinctive feature of foundation models is emergence
- Emergence means that the behavior of a system is implicitly induced rather than explicitly constructed

Foundation Models: Emergence

- A nice emergence example would be the GPT-3 model ([T. B. Brown et al. 2020](#)) composed of 175 billion parameters
- GPT-3 permits *in-context learning*, in which the language model can be adapted to a downstream task simply by providing it with a prompt
- This ability was not specifically trained for or anticipated to arise

Foundation Models: In-Context Learning

Let's make a small in-context learning experiment. Open ChatGPT and type:

```
review: delicious food ; sentiment: positive  
review: the food is awful ; sentiment: negative  
review: Terrible dishes! ; sentiment: negative  
review: Good meal! ; sentiment:
```

Why Foundation Models?

- The term describes a completely new trend in AI that is not captured by other terms
- This new trend is characterized by their sociological impact and broad shift they have introduced in AI research
- The word *foundation* alludes to the fact that these new models are incomplete but they constitute the basis for many task-specific models through adaptation
- The word also tries to capture the impact of the model quality on its many potential applications: is the model well or poorly constructed?

- Foundation models have broad sociological ramifications:
 - potential exacerbation of social inequities
 - economic impact due to increased capabilities
 - environmental impact due to increased computational demands
 - concerns about misuse, legal and ethics issues

The Future of Foundation Models

- There exist significant financial motivations to expand the abilities and magnitude of foundational models
- A continuous advancement in technology in the upcoming years can be expected
- Since foundation models rely on emergent behavior, its widespread deployment should be made with caution
- Foundation models appeared almost exclusively in industry but will benefit from the guidance of academia
- The negative trend that the most powerful and latest models are no longer publicly released constitutes an important problem



Amengual, Juan C., José M. Benedí, Francisco Casacuberta, Asunción Castaño, Antonio Castellanos, Víctor M. Jiménez, David Llorens, Andrés Marzal, Moisés Pastor, Federico Prat, Enrique Vidal, and Juan M. Vilar (2001). **“The EuTrans-I Speech Translation System”**. In: *Machine Translation* 15.2, pp. 75–103. DOI: 10.1023/A:1011116115948.



Bahdanau, Dzmitry, Kyung Hyun Cho, and Yoshua Bengio (Jan. 2015). **“Neural machine translation by jointly learning to align and translate”**. In: 3rd International Conference on Learning Representations, ICLR 2015 ; Conference date: 07-05-2015 Through 09-05-2015.



Bommasani, Rishi et al. (2021). **“On the Opportunities and Risks of Foundation Models”**. In: *ArXiv*. URL: <https://crfm.stanford.edu/assets/report.pdf>.



Brown, Peter F., Stephen A. Della Pietra, Vincent J. Della Pietra, and Robert L. Mercer (1993). “**The Mathematics of Statistical Machine Translation: Parameter Estimation**”. In: *Computational Linguistics* 19.2, pp. 263–311.



Brown, Tom B, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. (2020). “**Language models are few-shot learners**”. In: *arXiv preprint arXiv:2005.14165*.



Cho, Kyunghyun, Bart van Merriënboer, Dzmitry Bahdanau, and Yoshua Bengio (Oct. 2014). “**On the Properties of Neural Machine Translation: Encoder–Decoder Approaches**”. In: *Proceedings of SSST-8, Eighth Workshop on Syntax, Semantics and Structure in Statistical Translation*. Doha, Qatar: Association for Computational Linguistics, pp. 103–111.



Dauphin, Yann N., Angela Fan, Michael Auli, and David Grangier (2016). “**Language Modeling with Gated Convolutional Networks.**”. In: *CoRR* abs/1612.08083. URL: <http://dblp.uni-trier.de/db/journals/corr/corr1612.html#DauphinFAG16>.



Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2019). “**BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding**”. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*.



Hochreiter, Sepp and Jürgen Schmidhuber (1997). “**Long Short-Term Memory**”. In: *Neural Computation* 9.8, pp. 1735–1780.



Kim, Yoon (Oct. 2014). “**Convolutional Neural Networks for Sentence Classification**”. In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Doha, Qatar: Association for Computational Linguistics, pp. 1746–1751. DOI: 10.3115/v1/D14-1181. URL: <https://aclanthology.org/D14-1181>.



Koehn, Philipp (2005). “**Europarl: A Parallel Corpus for Statistical Machine Translation**”. In: *Proceedings of Machine Translation Summit X*. Phuket, Thailand, pp. 79–86. URL: <https://aclanthology.org/2005.mtsummit-papers.11/>.



Mikolov, Tomas, Martin Karafiát, Lukás Burget, Jan Cernocký, and Sanjeev Khudanpur (2010). “**Recurrent neural network based language model**”. In: *INTERSPEECH*. Ed. by Takao Kobayashi, Keikichi Hirose, and Satoshi Nakamura. ISCA, pp. 1045–1048. URL: <http://dblp.uni-trier.de/db/conf/interspeech/interspeech2010.html#MikolovKBCK10>.



Radford, Alec, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever (2018). “**Improving language understanding by generative pre-training**”. In.



Sutskever, Ilya, Oriol Vinyals, and Quoc V. Le (2014). “**Sequence to Sequence Learning with Neural Networks**”. In: *Proceedings of the 27th International Conference on Neural Information Processing Systems*. NIPS’14. Montreal, Canada: MIT Press, pp. 3104–3112.



Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin (2017). “**Attention is all you need**”. In: *Advances in Neural Information Processing Systems*, pp. 5998–6008.



Voita, Elena, David Talbot, Fedor Moiseev, Rico Sennrich, and Ivan Titov (July 2019). **“Analyzing Multi-Head Self-Attention: Specialized Heads Do the Heavy Lifting, the Rest Can Be Pruned”**. In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*. Florence, Italy: Association for Computational Linguistics, pp. 5797–5808. DOI: 10.18653/v1/P19-1580. URL: <https://aclanthology.org/P19-1580>.